

# COMP 206 – Software Systems

Lecture 22 – November 28, 2018


Coordination of multi process programs

# Last time

- Our processes' access to the CPU is through the scheduler
- We cannot guarantee what order jobs will run in, which can lead to bad outcomes when sharing a task:
  - "add\_ten.bash" actually sets the file back to 1!
  - The Dining Philosophers can end up starving unless they have a good plan


# Dining Philosophers again

- Why did we play this goofy game?
  - Philosophers are processes
  - Sticks are resources (a file, a hardware component, etc)
  - Eating means being able to do our desire work: read from a camera and post our selfie online
- The solutions to Dining Philosophers have a relation to coordination mechanisms in software systems



# Dining Philosophers solutions

- The key is to avoid the "one stick each" situation, and any equivalent state (such as alternating pick-up/put-down forever)
- We need a "tie-breaker". Lets you guarantee someone will get both sticks.
  - Recall: Numbering chopsticks works
  - Recall: Calling a "waiter" also works
  - Break the rules: communication would also work

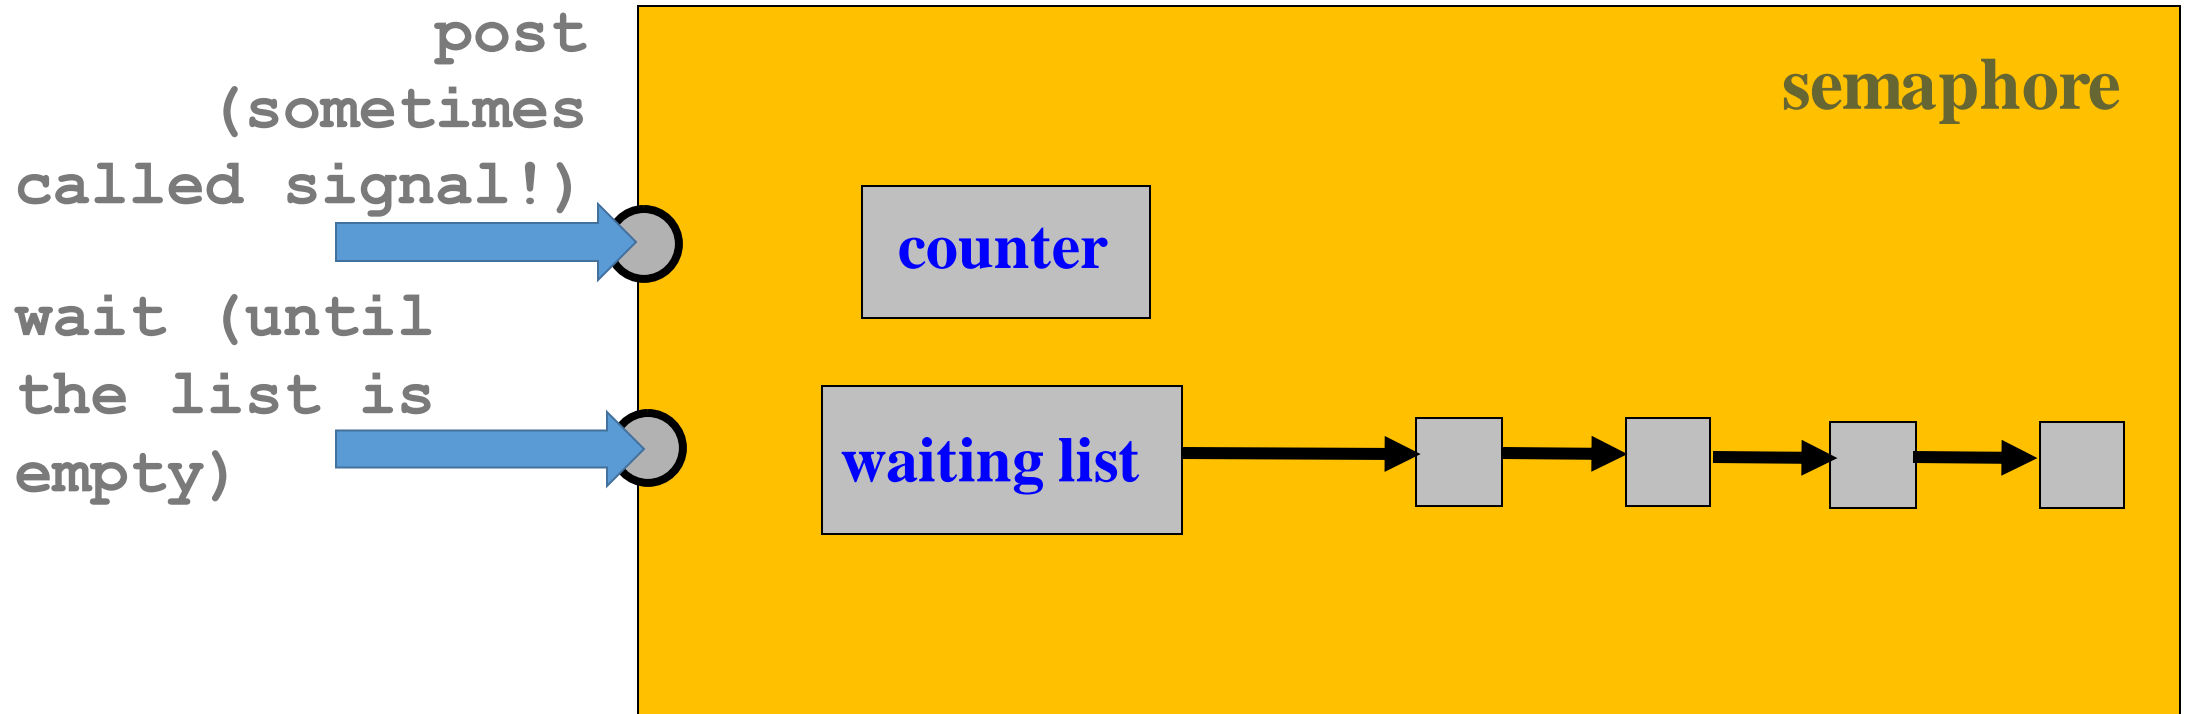


# Tie-breaking with semaphores

- Idea: have a central agent (the operating system) keep track of who asked first and only let one process execute at a time
- The process that asks first gets to proceed right away, and must then also say when it's finished
- All other processes are blocked by the central agent, and allowed to proceed one at a time.

# Semaphores

- A *semaphore* is an object that consists of a **counter**, a **waiting list** of processes and two **methods** (e.g., functions): `post` and `wait`.



# Semaphore Method: wait

```
void wait(sem S) {  
    S.count--;  
    if (S.count < 0) {  
        add the caller to the waiting list ;  
        block() ;  
    }  
}
```

- ❑ After decreasing the counter by 1, if the counter value becomes negative, then
  - ❖ add the caller to the waiting list, and then
  - ❖ block the caller.
  - ❖ if all processes call wait(), only "the initial value of count" can run at once!

# Semaphore Method: post

```
void post(sem S) {  
    S.count++;  
    if (S.count <= 0) {  
        remove a process P from the waiting list;  
        resume( P) ;  
    }  
}
```

- ❑ After increasing the counter by 1, if the new counter value is not positive, then
  - ❖ remove a process **P** from the waiting list,
  - ❖ resume the execution of process **P**, and return
  - ❖ recall wait() let only "initial count" jobs start. Post means "I finished", the next in line can start.



wait

## Important Note: 1/4

post

```
S.count--;  
if (S.count<0) {  
    add to list ;  
    block() ;  
}
```

```
S.count++;  
if (S.count<=0) {  
    remove P ;  
    resume(P) ;  
}
```

- ❑ If  $S.count < 0$ ,  $abs(S.count)$  is the number of waiting processes.
- ❑ This is because processes are added to (*resp.*, removed from) the waiting list only if the counter value is  $< 0$  (*resp.*,  $\leq 0$ ).

wait

## Important Note: 2/4

post

```
S.count--;  
if (S.count<0) {  
    add to list ;  
    block() ;  
}
```

```
S.count++;  
if (S.count<=0) {  
    remove    P ;  
    resume (   P) ;  
}
```

- ❑ The waiting list can be implemented with a queue if FIFO order is desired.
- ❑ However, the correctness of a program should not depend on a particular implementation of the waiting list.
- ❑ Your program should not make any assumption about the ordering of the waiting list.

wait

# Important Note: 3/4 post

```
S.count--;  
if (S.count<0) {  
    add to list ;  
    block() ;  
}
```

```
S.count++;  
if (S.count<=0) {  
    remove P ;  
    resume(P) ;  
}
```

- ❑ The caller may be blocked in the call to `wait()`.
- ❑ The caller never blocks in the call to `post()`.  
If `S.count > 0`, `post()` returns and the caller continues. Otherwise, a waiting process is released and the caller continues. In this case, *two* processes continue.

# The Most Important Note: 4/4

wait

```
S.count--;  
if (S.count<0) {  
    add to list ;  
    block() ;  
}
```

post

```
S.count++;  
if (S.count<=0) {  
    remove    P ;  
    resume (   P) ;  
}
```

- ❑ `wait()` and `post()` must be executed *atomically* (i.e., as one **uninterruptible** unit).
- ❑ Otherwise, *race conditions* may occur.
- ❑ **Exercise:** use execution sequences to show race conditions if `wait()` and/or `post()` is not executed atomically.

# Three Typical Uses of Semaphores

□ There are three typical uses of semaphores:

❖ **mutual exclusion:**

Mutex (*i.e.*, *Mutual Ex*clusion) locks

❖ **count-down lock:**

Keep in mind that semaphores have a counter.

❖ **notification:**

Indicate an event has occurred.

# Use 1: Mutual Exclusion (Lock)

*initialization is important*

```
sem S = sem_init(1);  
int    count = 0;
```

## Process 1

```
while (1) {  
    // do something  
    wait(S);  
    count++;  
    post(S);  
    // do something  
}
```

## Process 2

```
while (1) {  
    entry // do something  
    wait(S);  
    count--;  
    post(S);  
    exit // do something  
}
```

- ❑ **Question:** What if the initial value of **S** is zero?
- ❑ **S** is a *binary semaphore* (similar to a *lock*).

## Use 2: Count-Down Counter

```
sem S = sem_init(3);
```

### Process 1

```
while (1) {  
    // do something  
    wait(S);  
    at most 3 processes can be here!!!  
    post(S);  
    // do something  
}
```

### Process 2

```
while (1) {  
    // do something  
    wait(S);  
    post(S);  
    // do something  
}
```

- After **three** processes pass through `wait()`, this section is locked until a process calls `post()` .

# Use 3: Notification

```
sem S1 = sem_init(1), S2 = sem_init(0);


      process 1                                process 2
while (1) {                                while (1) {
    // do something                          // do something
    wait(S1);      notify                   wait(S2);
    printf("1");
    post(S2);      notify                   post(S1);
    // do something
}                                           }
```

- ❑ Process 1 uses `post(S1)` to notify process 2, indicating “**I am done. Please go ahead.**”
- ❑ The output is 1 2 1 2 1 2 .....
- ❑ What if both S1 and S2 are both 0's or both 1's?
- ❑ What if S1 = 0 and S2 = 1?



# Dining Philosophers With Semaphors

- We have the tool now!
- Can you think of how to apply the `wait()` and `signal()` operations of a semaphore to ensure the Dining Philosophers can eat in peace?
  - How many semaphores to use?
  - What code to run on each Philosopher?



# Dining Philosopher: Ideas

- ❑ Chopsticks are shared items (by two philosophers) and must be protected.
- ❑ Each chopstick has a semaphore with initial value 1.
- ❑ A philosopher calls `wait()` before they pick up a chopstick and calls `post()` to release it.

```
sem C[5];  
// init all to 1;
```

*outer critical section*

*inner critical*

`wait(C[i]);` *section*

`wait(C[(i+1)%5]);`

**has 2 chops and eats**

`post(C[(i+1)%5]);`

`post(C[i]);`

# Dining Philosophers: Proposal

```
sem C[5]; // init all to 1;
```

For all philosophers  $i$  (*in parallel*)

```
while (1) {  
    // thinking  
    wait (C[i]);  
    wait (C[(i+1)%5]);  
    // eating  
    post (C[(i+1)%5]);  
    post (C[i]);  
    // finishes eating  
}
```

*wait for my left chop*

*wait for my right chop*


*release my right chop*

*release my left chop*

**Does this solution work?**

# Dining Philosophers: Deadlock!

- No! This was the first solution we saw that didn't work.
- The semaphors here ensure the consistency of the game: no chopstick shared
- But, the coordination *protocol* ends in *deadlock*



# Dining Philosophers: A Better Idea

```
sem C[5]; // init all to 1;
```

**philosopher  $i$  (0, 1, 2, 3)**

```
while (1) {  
    // thinking  
    wait(C[i]);  
    wait(C[(i+1)%5]);  
    // eating  
    post(C[(i+1)%5]);  
    post(C[i]);  
    // finishes eating;  
}
```

*lock left chop first*

**Philosopher 4: the weirdo**

```
while (1) {  
    // thinking  
    wait(C[(i+1)%5]);  
    wait(C[i]);  
    // eating  
    post(C[i]);  
    post(C[(i+1)%5]);  
    // finishes eating  
}
```

*lock right chop first*

# Dining Philosophers: Questions

- The following are some important questions to think over.
  - ❖ We choose philosopher 4 to be the weirdo. Does this choice matter?
  - ❖ Show that this solution does not cause *circular waiting*.
  - ❖ Show that this solution will not have *circular waiting* if we have more than 1 and less than 5 weirdoes.

# Semaphores in C/Linux

- Functionality included in:
  - `#include <semaphore.h>`
  - `#include <fcntl.h>`
- Sits on top of the system call ***semctl()***
  - We will not look at that directly, only the easier C interface
- Allows the use of `wait()` and `post()` between processes and Linux. These tools and some experience give us full control over process ordering!
- More info: [http://man7.org/linux/man-pages/man7/sem\\_overview.7.html](http://man7.org/linux/man-pages/man7/sem_overview.7.html)

# Synchronized counting with semaphores

- Recall our fork() example from last class that counted to 200. The two processes did not stay in synch. Can you think of how to synchronize them so we see: "1, 1, 2, 2, 3, 3, .... , 199, 199, 200, 200" each time?

```
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
#define MAX_COUNT 200
#define BUF_SIZE 100
void main(void)
{
    pid_t pid;
    int i;
    char buf[BUF_SIZE];

    fork();
    pid = getpid();
    for (i = 1; i <= MAX_COUNT; i++) {
        sprintf(buf,
                "This line is from pid %d, value = %d\n", pid, i);
        write(1, buf, strlen(buf));
    }
}
```



# Solution:

```
1 #include <stdio.h>
2 #include <string.h>
3 #include <unistd.h>
4
5 #include <semaphore.h>
6 #include <fcntl.h>
7
8 #define MAX_COUNT 200
9 #define BUF_SIZE 100
10
11 sem_t *s1;
12 sem_t *s2;
13
14 void main(void)
15 {
16     pid_t pid;
17     int i;
18     char buf[BUF_SIZE];
19
20     s1 = sem_open( "first", O_CREAT, 0666, 1 );
21     s2 = sem_open( "second", O_CREAT, 0666, 0 );
22
23     int x = fork();
24     pid = getpid();
25     for (i = 1; i <= MAX_COUNT; i++) {
26
27         if( x == 0 ){
28             sem_wait( s1);
29         }
30         else{
31             sem_wait( s2 );
32         }
33
34         sprintf(buf, "This line is from pid %d, value = %d\n", pid, i);
35         write(1, buf, strlen(buf));
36
37         if( x == 0 ){
38             sem_post( s2);
39         }
40         else{
41             sem_post( s1 );
42         }
43     }
44 }
45
```

This line is from pid 22918, value = 1

This line is from pid 22917, value = 1

This line is from pid 22918, value = 2

This line is from pid 22917, value = 2

...

This line is from pid 22918, value = 199

This line is from pid 22917, value = 199

This line is from pid 22918, value = 200

This line is from pid 22917, value = 200

# Solution:

- This is the "notification" pattern:
  - Each side has to wait until the other has posted to advance
  - We start the semaphores at 0 and 1 to specify explicitly who gets to run first
  - Afterwards, the order is not guaranteed, but the fact that they stay within 1 number is

```
1 #include <stdio.h>
2 #include <string.h>
3 #include <unistd.h>
4
5 #include <semaphore.h>
6 #include <fcntl.h>
7
8 #define MAX_COUNT 200
9 #define BUF_SIZE 100
10
11 sem_t *s1;
12 sem_t *s2;
13
14 void main(void)
15 {
16     pid_t pid;
17     int i;
18     char buf[BUF_SIZE];
19
20     s1 = sem_open( "first", O_CREAT, 0666, 1 );
21     s2 = sem_open( "second", O_CREAT, 0666, 0 );
22
23     int x = fork();
24     pid = getpid();
25     for (i = 1; i <= MAX_COUNT; i++) {
26
27         if( x == 0 ){
28             sem_wait( s1 );
29         }
30         else{
31             sem_wait( s2 );
32         }
33
34         sprintf(buf, "This line is from pid %d, value = %d\n", pid, i);
35         write(1, buf, strlen(buf));
36
37         if( x == 0 ){
38             sem_post( s2 );
39         }
40         else{
41             sem_post( s1 );
42         }
43     }
44 }
45 |
```

# Further thoughts

- Programming languages implement many additional features to allow mutual exclusion, coordination etc.
- We will sometimes see Linux semaphores directly in your C code, but more often they will be hidden within other features :
  - Shared memory (exclusion tied directly to memory access)
  - Mutex (implements only the binary lock for critical sections)
  - Scoped lock (no need to call init or signal, the variable's lifetime on the stack controls the exclusion)
- However, the concepts of all these are the same and now we've seen under the hood a bit. Let's think back on the key elements.

# Coordination Summary Points

- Processes must **actively synchronize** due to unpredictable scheduling
- **Critical sections** are portions of shared/parallel code where it is essential only a single process should be operating at a time:
  - File reading/writing, updating memory locations, eating spaghetti
- A **race condition** is a situation where performance depends on the ordering of processes. We must avoid these!
- Semaphores are one coordination construct: wait(), post(), a counter and a list of waiting processes
- The key to implementing semaphores (and all other coordination) is that all calls are **atomic operations**. This means behavior is guaranteed even if 2 processes call the function at exactly the same time (or close). Atomic operations **break ties!**